

1. Why do we use Kubernetes? what is a container orchestration tool?

Answer: container orchestration tool which is used to manage the containers by automating the various concerns around container “running.”

Example of container orchestration tool-

Kubernetes

OpenShift

Docker Swarm

We use Kubernetes because of these features-

- Kubernetes is an open-source tool
- Automation of deployment and scalability
- It provides load balancing
- it has a feature of self healing
- it makes the application highly available
- it uses the resources efficiently
- you can make n number of containers through one command.
- Kubernetes works on both vertical and horizontal scaling.
- Works on client server model
- Kubernetes support

2. What were the challenges we faced in docker? Mention the answer in bullet points

1. There has no Autoscaling feature
2. There has no Load balancing feature
3. There has no Self – healing feature
4. you can not make multiple container through on command.
You will have to run command manually for each container.
5. Docker prefer vertical scaling.

3. What components are used in Kubernetes architecture and what model do we use to manage Kubernetes on our system?

Create Kubernetes cluster (take whichever machine you like atleast use 1 master node and 2 worker nodes)

Answer- In Kubernetes architecture there is cluster which is a combination of master node and worker node.

Master node/control panel component-

- kube-apiserver
- etcd
- kube-scheduler
- kube-controller-manager

Worker node component-

- kubelet
- kube-proxy
- Container runtime

```

root@master:~# kubectl apply -f calico.yaml
poddisruptionbudget.policy/calico-kube-controllers configured
serviceaccount/calico-kube-controllers unchanged
serviceaccount/calico-node unchanged
configmap/calico-config unchanged
customresourcedefinition.apiextensions.k8s.io/bgpconfigurations.crd.projectcalico.org configured
customresourcedefinition.apiextensions.k8s.io/bgppeers.crd.projectcalico.org configured
customresourcedefinition.apiextensions.k8s.io/blockaffinities.crd.projectcalico.org configured
customresourcedefinition.apiextensions.k8s.io/caliconodestatuses.crd.projectcalico.org configured
customresourcedefinition.apiextensions.k8s.io/clusterinformations.crd.projectcalico.org configured
customresourcedefinition.apiextensions.k8s.io/felixconfigurations.crd.projectcalico.org configured
customresourcedefinition.apiextensions.k8s.io/globalnetworkpolicies.crd.projectcalico.org configured
customresourcedefinition.apiextensions.k8s.io/globalnetworksets.crd.projectcalico.org configured
customresourcedefinition.apiextensions.k8s.io/hostendpoints.crd.projectcalico.org configured
customresourcedefinition.apiextensions.k8s.io/ipamblocks.crd.projectcalico.org configured
customresourcedefinition.apiextensions.k8s.io/ipamconfigs.crd.projectcalico.org configured
customresourcedefinition.apiextensions.k8s.io/ipamhandles.crd.projectcalico.org configured
customresourcedefinition.apiextensions.k8s.io/ippools.crd.projectcalico.org configured
customresourcedefinition.apiextensions.k8s.io/ipreservations.crd.projectcalico.org configured
customresourcedefinition.apiextensions.k8s.io/kubecontrollersconfigurations.crd.projectcalico.org configured
customresourcedefinition.apiextensions.k8s.io/networkpolicies.crd.projectcalico.org configured
customresourcedefinition.apiextensions.k8s.io/networksets.crd.projectcalico.org configured
clusterrole.rbac.authorization.k8s.io/calico-kube-controllers unchanged
clusterrole.rbac.authorization.k8s.io/calico-node unchanged
clusterrolebinding.rbac.authorization.k8s.io/calico-kube-controllers unchanged
clusterrolebinding.rbac.authorization.k8s.io/calico-node unchanged
daemonset.apps/calico-node configured
deployment.apps/calico-kube-controllers unchanged
root@master:~# kubectl get nodes

```

NAME	STATUS	ROLES	AGE	VERSION
master.jyoti.com	Ready	control-plane	7m29s	v1.28.2
worker1.jyoti.com	Ready	<none>	6m42s	v1.28.2
worker2.jyoti.com	Ready	<none>	5m53s	v1.28.2

```

root@master:~#

```

4. What is a Pod? create a pod name grras which uses the image of nginx with port 80 allowed in the manifest file.

Answer- A pod is the smallest execution unit which is managed by Kubernetes.

```

root@master:~# vim pod.yaml
root@master:~# kubectl apply -f pod.yaml
pod/grras created
root@master:~# kubectl get pods

```

NAME	READY	STATUS	RESTARTS	AGE
grras	1/1	Running	0	18s

```

root@master:~# kubectl get pods -o wide

```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
grras	1/1	Running	0	34s	192.168.59.1	worker1.jyoti.com	<none>	<none>

```

root@master:~# cat pod.yaml
apiVersion: v1
kind: Pod
metadata:
  name: grras
spec:
  containers:
    - name: container1
      image: nginx:latest
      ports:
        - containerPort: 80

```

5. What are the strategies we use to achieve HA (high availability) for an application? (hint: Deployment strategies).

Answer:

1. Replication Controller

A *Replication Controller* ensures that a specified number of pod replicas are running at any one time.

Replication Controller makes sure that a pod or a homogeneous set of pods is always up and available.

2. Replica set

A Replica Set's purpose is to maintain a stable set of replica Pods running at any given time.

it is often used to guarantee the availability of a specified number of identical Pods

3. Deployment

A *Deployment* provides declarative updates for Pods and Replica Sets.

6. Create an RC name engineer (it should be run in a namespace called aaramb1, rc should manage 9 replicas), create one file called manual.txt, and write a command to scale up your pod replicas to 15 without changing it in the manifest file).

```
root@master:~# kubectl get namespace
NAME                STATUS   AGE
default             Active   4m16s
kube-node-lease     Active   4m16s
kube-public         Active   4m16s
kube-system         Active   4m16s
root@master:~# vim namespace.yaml
root@master:~# kubectl create -f namespace.yaml
namespace/aaramb1 created
root@master:~# kubectl get namespace
NAME                STATUS   AGE
aaramb1             Active   5s
default             Active   6m11s
kube-node-lease     Active   6m11s
kube-public         Active   6m11s
kube-system         Active   6m11s
root@master:~# ls
calico.yaml  namespace.yaml  snap
root@master:~# vim manual.txt
root@master:~# ls
calico.yaml  manual.txt  namespace.yaml  snap
root@master:~# vim rc.yaml
root@master:~# vim rc.yaml
root@master:~# kubectl create -f rc.yaml
replicationcontroller/engineer created
root@master:~# kubectl get rc
No resources found in default namespace.
root@master:~# kubectl get rc --namespace=aaramb1
NAME          DESIRED   CURRENT   READY   AGE
engineer      9         9         9       30s
root@master:~# kubectl describe rc engineer --namespace=aaramb1
error: unknown command "describe" for "kubectl"

Did you mean this?
  describe
root@master:~# kubectl describe rc engineer --namespace=aaramb1
Name:          engineer
Namespace:     aaramb1
Selector:      app=nginx
Labels:        app=nginx
Annotations:   <none>
Replicas:      9 current / 9 desired
```

Normal successful create 0/1 Replication controller created po

```
root@master:~# vim manual.txt
root@master:~# ./manual.txt
-bash: ./manual.txt: Permission denied
root@master:~# ls -l manual.txt
-rw-r--r-- 1 root root 72 Sep 28 05:41 manual.txt
root@master:~# chmod -R 777 manual.txt
root@master:~# ls -l manual.txt
-rwxrwxrwx 1 root root 72 Sep 28 05:41 manual.txt
root@master:~# ./manual.txt
replicationcontroller/engineer scaled
root@master:~# kubectl get rc --namespace=aaramb1
NAME          DESIRED   CURRENT   READY   AGE
engineer      15        15        15      6m30s
root@master:~# cat rc.yaml
apiVersion: v1
kind: ReplicationController
metadata:
  name: engineer
  namespace: aaramb1
spec:
  replicas: 9
  template:
    metadata:
      name: mypod
      labels:
        app: nginx
    spec:
      containers:
        - name: container1
          image: nginx:latest
          ports:
            - containerPort: 8
```

7. What is the difference between RC and RS? What do labels and selectors explain?

Replication Controller	Replica Set
The Replication Controller uses equality-based selectors to manage the pods.	ReplicaSets Controller uses set-based selectors to manage the pods.
The rolling-update command works with Replication Controllers	The rolling-update command won't work with ReplicaSets.
Replica Controller is deprecated and replaced by ReplicaSets.	Deployments are recommended over ReplicaSets.
Selector field is not mandatory in replication controller	Selector field is mandatory in replication controller

Labels : *Labels* are key/value pairs that are attached to objects such as Pods.

Labels can be used to organize and to select subsets of objects.

Labels can be attached to objects at creation time and subsequently added and modified at any time.

Each object can have a set of key/value labels defined. Each Key must be unique for a given object.

labels do not provide uniqueness. In general, we expect many objects to carry the same label(s).

Selectors: the client/user can identify a set of objects.

The label selector is the core grouping primitive in Kubernetes.

The API currently supports two types of selectors: *equality-based* and *set-based*.

8. Create a deployment in your system:

- The deployment name should be website
- Labels it uses must be app=website and website=app
- Replicas it should use 12
- Use the image nginx:latest for deployment

```
root@master:~# kubectl apply -f deploy.yaml
deployment.apps/website created
root@master:~# vim deploy.yaml
root@master:~# vim deploy.yaml
root@master:~# cat deploy.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: website
  labels:
    app: website
spec:
  replicas: 12
  selector:
    matchLabels:
      app: website
      website: app
  template:
    metadata:
      labels:
        app: website
        website: app
    spec:
      containers:
      - name: nginx
        image: nginx:latest
        ports:
        - containerPort: 80

root@master:~# kubectl get deployment
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
website       12/12   12           12          2m35s
root@master:~#
```



```

command terminated with exit code 1
root@master:~# kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
website-db46df88-4fzrg              1/1     Running   0           39m
website-db46df88-5n7k6              1/1     Running   0           39m
website-db46df88-85kvd              1/1     Running   0           39m
website-db46df88-bdgfg              1/1     Running   0           39m
website-db46df88-f6f2b              1/1     Running   0           39m
website-db46df88-gnqk8              1/1     Running   0           39m
website-db46df88-hkphl              1/1     Running   0           39m
website-db46df88-n2f2v              1/1     Running   0           39m
website-db46df88-ntk5b              1/1     Running   0           39m
website-db46df88-v2qm8              1/1     Running   0           39m
website-db46df88-vdpv9              1/1     Running   0           39m
website-db46df88-vvdj7              1/1     Running   0           39m
root@master:~# kubectl exec website-db46df88-5n7k6 -- curl http://localhost
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total   Spent    Left   Speed
  0     0    0     0    0     0      0      0  --:--:-- --:--:-- --:--:--    0<h1>i have completed my project</h1>
100   37 100   37    0     0  31951    0  --:--:-- --:--:-- --:--:--  37000
root@master:~#

```